



PRODUCT REQUIREMENTS DOCUMENT

AAA Insights

Collect surveys, reviews, ratings, and comments — then let AI turn them into customer, competitive, and relationship insight.

Version: v3 – Draft (Trust extension)
Date: 2026-07-14
Prepared by: Active AI Advisors
Prepared for: Paul Jamieson
Methodology: Grounded AI™ – Phase 1 (SPECIFY)

Applied AI for operators. · activeaiadvisors.com · paul@activeaiadvisors.com

WHAT CHANGED IN V3

Revised 2026-07-14 — Trust extension

- Brand Trust added as a first-class indicator that complements Brand Love (love is the emotional pull; trust is confidence the brand is reliable, honest, and acts in your interest).
- A Trust question type (single-item + optional driver battery: reliability, integrity, benevolence, security/privacy), a Trust Index, and trust inferred from open text.
- A Love × Trust segmentation — Devoted / Infatuated (fragile) / Dependable / At-risk — with a recommended action per quadrant.
- Grounded in the library: Wardani & Gustia (2016) — trust bridges satisfaction to attachment; Nobre (2011) — passion-love without trust is fragile.
(v2 added: aggregate unified-customer view, competitors, lawful collection under DPS-7, the Brand Love scale, and the data model.)

CONTENTS

- 01 Purpose
- 02 Target users & buyer
- 03 Product principles
- 04 Inputs
- 05 Outputs
- 06 Requirements (behaviors)
- 07 Invariants
- 08 Edge cases
- 09 Exclusions
- 10 Data, privacy & security
- 11 Non-functional requirements
- 12 Data model (logical)
- 13 Scope & phased roadmap
- 14 Open questions
- Approval

SECTION 01

Purpose

What the software is for — and who it serves.

SECTION 02

Target users & buyer

Two audiences, one product: credible analysis for the business, effortless response for the customer.

User role	What they need from AAA Insights
Owner / founder / GM (buyer)	A fast, trustworthy read on what customers think, how the brand compares to rivals, and what to fix first.
Marketing / brand lead	Themes, sentiment, brand-love, trust, quotes, and competitive position to guide messaging and positioning.
CX / customer-success manager	Early warning on dissatisfaction, eroding trust, drivers of churn, and a way to close the loop.
Product manager	Prioritized signal on what customers want changed, backed by evidence and volume.
Respondent (the customer)	A short, respectful, mobile-friendly way to give feedback — including a conversational option that feels heard.

SECTION 03

Product principles

Seven commitments that shape every design decision.

SECTION 04

Inputs

What the system receives.

- **I-1** Survey definition — built by an admin or AI-generated. Question types include single-select, multi-select, rating/scale, open text, the Brand Love scale (Love / Like / Ambiguity / Dislike / Hate), and the Trust battery (a single-item trust rating and/or a short driver battery: reliability, integrity, benevolence, security/privacy).
- **I-2** Survey response — ratings (incl. Brand Love and Trust), multiple-choice, and open-text comments. May be partial.
- **I-3** Conversational response — free-text turns in an AI-led adaptive interview, as a transcript.
- **I-4** Imported feedback — external reviews/ratings/comments for the company's own brand via CSV.
- **I-5** Respondent metadata — optional, non-identifying attributes (segment, channel, product, region, language).
- **I-6** Distribution request — publish a survey via link, email list, or widget.
- **I-7** Analysis query — a natural-language question asked of the feedback data.
- **I-8** Competitor configuration — the competitor brands to track: names, aliases, products, public sources.
- **I-9** Externally-collected feedback — reviews/ratings/comments about the own brand and tracked competitors from public web pages, review-site APIs, or a licensed provider (source URL, capture date, brand, rating, text).

SECTION 05

Outputs

What the system produces.

- **O-1** Collected response store — every response persisted under a common schema, attributable to source, brand, and campaign.
- **O-2** Theme analysis — AI-generated themes with counts, quotes, and trend; filterable by brand.
- **O-3** Sentiment analysis — positive/neutral/negative with intensity, aggregated by theme, segment, source, brand, and time.
- **O-4** Ratings & metrics — average ratings, distribution, response/completion rates, NPS/CSAT.
- **O-5** Insight report — narrative of what's happening, why, and what to do, with ranked actions; can include the competitive picture and the Love × Trust read. Exportable.
- **O-6** Answer to an analysis query — grounded strictly in the account's own data, with citations.
- **O-7** Alerts — notifications when a monitored signal crosses a threshold.

- **0-8** Respondent-facing survey — the rendered survey or conversational interview.
- **0-9** Competitive benchmark — own brand vs. competitors on ratings, Brand Love, Trust, sentiment, and themes over a chosen period.
- **0-10** Per-competitor analysis — for a selected competitor, an aggregate view with traceability.
- **0-11** Brand Love read — the distribution across Love / Like / Ambiguity / Dislike / Hate and a Brand Love Index (share of Love+Like minus share of Dislike+Hate), per company and competitor, over time.
- **0-12** Trust read — a Trust Index (net trust) plus a driver breakdown (reliability, integrity, benevolence, security/privacy), per company and competitor, over time.

O-13 — Love × Trust segmentation

Respondents and segments are placed in one of four quadrants, each with a recommended action:

Quadrant	Meaning and recommended action
Devoted	High love / high trust — your advocates. Protect and activate them.
Infatuated (fragile)	High love / low trust — passion without a safety net. One stumble triggers churn; shore up reliability and transparency.
Dependable	Low love / high trust — loyal by reliability, not emotion. Steady; deepen the relationship.
At-risk	Low love / low trust — churn risk and detractors. Intervene or triage.

SECTION 06

Requirements (behaviors)

Each is a single, testable behavior. IDs are cited by the Phase-3 test suite.

Survey creation & distribution

- **R-1** Create a survey with at least: single-select, multi-select, rating/scale, open text, the Brand Love scale, and the Trust battery (single-item and/or driver items).
- **R-2** Generate a draft survey from a plain-language objective; editable before sending.
- **R-3** Conditional logic: a question or branch is shown or skipped based on a previous answer.
- **R-4** Distribute by (a) link, (b) email list, and (c) embeddable widget. (Email is Phase-1; link + widget are MVP.)
- **R-5** A respondent completes a survey on mobile and desktop without creating an account.

- **R-6** Partial responses are recorded and marked incomplete, not discarded.

Conversational (AI-led) surveys

- **R-7** Enable a conversational mode where an AI interviewer asks questions and relevant follow-up probes.
- **R-8** The interviewer stays on the admin-defined objective and does not ask outside the topic scope.
- **R-9** The interviewer ends after a configurable max number of exchanges or when the objective is met.
- **R-10** Every transcript is stored and analyzable by the same engine as structured responses.

Unified feedback hub

- **R-11** Import external feedback via CSV, mapping columns to the unified schema (brand, source, date, rating, text, segment).
- **R-12** All feedback is queryable together, filterable by brand, source, campaign, date, segment, rating, sentiment.
- **R-13** De-duplicate obviously identical items (same source + text + timestamp).

Brand Love, Trust & metrics

- **R-14** Produce theme analysis with counts and representative quotes for any filtered set.
- **R-15** Assign sentiment to each open-text response and aggregate across chosen dimensions.
- **R-16** For every theme, sentiment aggregate, Brand Love read, and Trust read, open the exact underlying responses (traceability — INV-3).
- **R-30** Record Brand Love on the five-point scale, normalize to an ordinal, and compute the Brand Love distribution and Index (O-11) — including inferring a read from open text where no direct rating exists, labeled inferred (INV-4).
- **R-31** Support a Trust question type — a single-item trust rating and an optional multi-item driver battery (reliability, integrity, benevolence, security/privacy).
- **R-32** Record trust; compute a Trust Index and per-driver breakdown (O-12) for any filtered set — including inferring a trust read from open text where no direct rating exists, labeled inferred (INV-4), the same way Brand Love is inferred (R-30).
- **R-33** Produce the Love × Trust segmentation (O-13) with a recommended action per quadrant, available in the insight report (O-5) and the competitive benchmark (O-9).

Unified customer & competitive insight

- **R-26** Produce an aggregate, cross-source "unified customer" view WITHOUT building or requiring an identity-linked individual profile (INV-9).
- **R-24** Define and manage a set of competitor brands to track (I-8).

- **R-25** Collect publicly available reviews/ratings/comments about the own brand and tracked competitors from configured sources, subject to §10 (DPS-7) and a legal-review gate before live collection is enabled.
- **R-27** Benchmark the company's brand vs. competitors on ratings, Brand Love, Trust, sentiment, and themes over a chosen period (O-9).
- **R-28** For a selected competitor, produce per-competitor aggregate analysis (O-10) with traceability.
- **R-29** All brand/competitor analysis is filterable by brand, source, date, segment, and sentiment, and every competitor figure links to its source items.

AI analysis, accounts, roles, export

- **R-17** Ask a natural-language question and receive an answer grounded only in the account's own data, with citations.
- **R-18** Generate an insight report: narrative, metrics, top themes, sentiment, Brand Love, Trust and the Love × Trust read, competitive comparison, quotes, ranked actions.
- **R-19** Compute response rate, completion rate, average rating, rating distribution, and NPS/CSAT when questions qualify.
- **R-20** Configure an alert on a signal (negative-sentiment share, average rating, Brand Love Index, Trust Index, new theme) with notification when crossed.
- **R-21** At least two roles: Owner/Admin (account, users, billing) and Member (builds surveys, views analysis).
- **R-22** Export any analysis view and the insight report (CSV for raw data; PDF and/or slides).
- **R-23** Delete a survey, campaign, or response, and export or delete all account data.

SECTION 07

Invariants

Conditions that must hold in every state of the system.

- **INV-1** Every response is attributable to exactly one source AND one brand. No response exists without a known origin.
- **INV-2** Counts are never negative and never exceed responses collected; percentages come from the current filter, each response counted once per dimension.
- **INV-3** Every theme, aggregate, Brand Love read, Trust read, insight, recommendation, and query answer is traceable to its underlying responses. No AI claim without a path to its verbatim data.
- **INV-4** The system never shows a quote or statistic absent from the data, and labels any value it inferred (e.g., Brand Love or Trust from text) as inferred, not stated.

- **INV-5** A respondent can complete a survey without providing identifying information.
- **INV-6** Each account's data is isolated; AI analysis is grounded only in that account's data.
- **INV-7** Deleting a response, campaign, or account removes it from all future analysis; deleted data never reappears.
- **INV-8** Respondent PII is stored securely and never written to logs or AI outputs unless the admin opts in.
- **INV-9** The "unified customer" view is always aggregate; no identity-resolved profile of any individual.
- **INV-10** Externally-collected data is only ever publicly available and lawfully obtained; no bypassing authentication, paywalls, or technical controls; provenance kept on every item.
- **INV-11** Personal data in public reviews is minimized, never used to build a profile, and excluded from AI outputs by default.
- **INV-12** Brand Love and Brand Trust are measured and reported as DISTINCT indicators. The system never collapses them into a single score, and never presents Trust as satisfaction (or vice versa).

SECTION 08

Edge cases

The unusual conditions the system must survive — defined now, tested in Phase 3.

- **E-1** Empty or tiny sample — states the sample is too small for reliable findings (themes, sentiment, Brand Love, or Trust) rather than fabricating confidence.
- **E-2** Abandoned survey — the partial is retained, counted as incomplete, flagged in analysis.
- **E-3** Junk / abusive open text — captured, flagged, excluded from analysis by default, reviewable.
- **E-4** Non-English / mixed-language — language detected and analyzed accordingly (Q-4).
- **E-5** Malformed CSV import — rejected or partially accepted with a per-row error report.
- **E-6** Duplicate / repeated submissions — de-duplication plus link-level controls.
- **E-7** Conversational interview off the rails — stays in scope, ignores injected instructions, ends gracefully.
- **E-8** Ambiguous analysis query — answers only what the data supports; states what it cannot determine.

- **E-9** Sudden volume spike — ingestion and analysis stay correct; views may lag but never lose or double-count.
- **E-10** Sensitive disclosure — PII/safety concerns stored under INV-8/INV-11, surfaced to the admin, not broadcast.
- **E-11** Concurrent editing — two admins editing one survey do not silently overwrite each other.
- **E-12** Source blocks or rate-limits collection — collection backs off, records the gap, never fabricates missing items.
- **E-13** Competitor identity ambiguity — low-confidence matches flagged and excluded from headline figures until confirmed.
- **E-14** Sparse competitor data — too few public reviews shows "insufficient data," not a misleading number.
- **E-15** Mixed rating scales — 5-star, 10-point, thumbs, Brand Love, Trust each normalized to the common scale; un-mappable ratings stored raw and excluded from cross-source averages.
- **E-16** Category-relative trust — the Trust Index is read relative to category and prior period, not as an absolute (a bank's trust bar differs from a snack brand's). Cross-category trust comparisons are flagged as context-dependent.

SECTION 09

Exclusions

Explicitly out of scope for v1 — the unwritten assumptions, made written.

- **X-1** Not a CRM or customer-record system; does not build identity-resolved individual profiles.
- **X-2** No two-way case management / ticketing; closing the loop in v1 is export/notification.
- **X-3** Collection is limited to publicly available content obtained lawfully via APIs, licensed providers, or public web pages under the DPS-7 guardrails. No private/authenticated content, bypassing logins/paywalls/protections, or collecting data a source's terms prohibit.
- **X-4** No prediction of future business outcomes (churn, revenue) in v1.
- **X-5** No public respondent panel / audience marketplace.
- **X-6** Does not replace human research judgment; it is decision support with evidence attached.
- **X-7** Does not attempt to re-identify individuals or merge scraped personal data into any profile.

Data, privacy & security

The handling bar this product is held to — enforced in build and verification.

- **DPS-1** The system is a data processor for the account's feedback; the company is the controller.
- **DPS-2** Respondent PII is minimized, encrypted at rest and in transit, access-controlled by role, and excluded from logs and AI outputs by default.
- **DPS-3** Data-subject requests are supported: export and deletion on request.
- **DPS-4** AI processing is isolated to the account; account data is not used to train shared/base models without opt-in.
- **DPS-5** A configurable retention period is supported for both first-party and collected data.
- **DPS-6** The full security checklist is applied during build and verification.
(Security/privacy is also a Trust driver measured in O-12 — how safe customers feel is part of the trust picture.)

DPS-7 — Web / data-collection compliance

Lawful and ethical by construction — with a legal-review gate

Collect only publicly accessible content. Respect robots.txt and reasonable rate limits.

Never bypass authentication, paywalls, or technical protections. Prefer official APIs and licensed providers over page collection.

Keep provenance (source, URL, capture date) on every item. Minimize and never profile personal data in reviews (INV-11). Honor source-specific terms.

The specific sources and methods are subject to the client's legal review and written sign-off BEFORE live collection is enabled — a gate, not a default.

Active AI Advisors is not providing legal advice; the client should confirm the approach with counsel.

Non-functional requirements

Speed, reliability, access, and honesty — stated as testable budgets.

- **NFR-1** Performance — a survey page loads under 2 s on mobile; a submission acknowledges under 1 s.

- **NFR-2** Analysis latency — up to the MVP size (~5,000 items, Q-6) returns within a stated budget; larger sets run as background jobs with progress.
- **NFR-3** Availability — the collection endpoint targets a stated uptime; ingestion is durable even if analysis is degraded.
- **NFR-4** Accessibility — respondent-facing surveys meet WCAG 2.1 AA.
- **NFR-5** Scalability path — grows from SMB volumes toward larger sets without a rewrite.
- **NFR-6** Explainability — every AI output carries the data behind it and a confidence signal, including inferred Brand Love and Trust reads.
- **NFR-7** Collection freshness & politeness — collected data shows a freshness/coverage indicator; collection runs at a polite, configurable rate and records gaps.

SECTION 12

Data model (logical)

The records the system keeps and how they relate. Physical schema is an Architecture-phase decision; this defines what is stored, not how.

Brand — *a company or competitor whose feedback is analyzed*

Field	Meaning
brand_id	Unique id
name / aliases	Display name and alternates used to match reviews
type	own or competitor
products	Optional product / line names
tracked	Whether active collection is enabled

Source — *where a piece of feedback came from*

Field	Meaning
source_id	Unique id
type	survey · conversational · import_csv · web · api · provider
name / url	Human name and, for web/API, the location
terms_note	Compliance note / permitted-use flag (DPS-7)

FeedbackRecord — *the core record — one rating and/or comment about one brand*

Field	Meaning
record_id	Unique id

Field	Meaning
account_id	Owning tenant (INV-6)
brand_id	Which brand this is about — own or competitor (INV-1)
source_id	Where it came from (INV-1)
captured_at / date	When the feedback was given or collected
rating_raw / rating_scale	Original rating and its scale (5_star, 10_pt, nps, csat, brand_love, trust, ...)
rating_norm	Rating normalized to a common 1–5 ordinal (E-15)
brand_love	Love / Like / Ambiguity / Dislike / Hate, if applicable — ordinal 5→1
trust (v3)	Trust rating if directly asked — normalized 1–5
trust_drivers (v3)	Optional per-driver scores/tags: reliability, integrity, benevolence, security_privacy
comment_text	The open-text comment
language	Detected language (E-4)
segment / region / channel	Optional non-identifying metadata (I-5)
is_complete	Complete vs. partial (R-6)
flags	junk · abuse · safety · low_confidence (E-3, E-13)
provenance	source_url + capture_date, or import_batch (INV-10)

Sentiment — *derived, one per record*

Field	Meaning
record_id	The record it describes
polarity / intensity	positive / neutral / negative, with strength
model_version / confidence	What produced it and how sure (NFR-6)

Trust (v3) — *derived, one per record where trust is present or inferred*

Field	Meaning
record_id	The record it describes
trust_score	Normalized trust value
drivers	reliability / integrity / benevolence / security_privacy sub-scores
model_version / confidence	What produced it and how sure (NFR-6)

Field	Meaning
inferred	True if derived from open text rather than a direct rating (INV-4)

MetricSnapshot — *a computed rollup for a brand over a period and filter (O-4, O-9, O-11, O-12)*

Field	Meaning
brand_id / period	Brand and time window
metric	avg_rating · nps · csat · brand_love_index · trust_index (v3) · trust_<driver> (v3) · neg_sentiment_share · response_rate · completion_rate
value / filter_context	The number and the filter it was computed under (INV-2)

Love × Trust segment (v3) & the "unified customer" view

Love × Trust segment — computed from the Brand Love and Trust metrics for a brand/segment; each respondent or segment falls in one quadrant (Devoted / Infatuated / Dependable / At-risk) with counts and shares (O-13).

Unified customer view — there is deliberately NO individual-Customer entity. It aggregates FeedbackRecords for an own brand across all sources and time, grouped by Segment — a population read, not a person read (R-26, INV-9).

SECTION 13

Scope & phased roadmap

Only Phase 0 is committed by this document. Each later phase re-enters the pipeline before it is built.

Phase 0 — MVP (the focused, buildable core)

- Survey builder: core types + Brand Love scale + Trust battery + conditional logic (R-1, R-3, R-30, R-31) • AI-drafted survey (R-2)
- Distribution by link + widget (R-4) • Mobile, account-free capture incl. partials (R-5, R-6)
- Conversational AI survey with scope control + turn limits (R-7–R-10) • CSV import into the unified hub (R-11, R-13)
- Unified cross-source analysis + aggregate unified-customer view (R-12, R-26)
- AI theme + sentiment + Brand Love + Trust with traceability (R-14–R-16, R-30, R-32)
- Core metrics incl. NPS/CSAT + Brand Love Index + Trust Index (R-19, O-11, O-12)
- Love × Trust segmentation with recommended actions (R-33, O-13)

- Insight report with ranked actions (R-18, R-22) • Two roles, account isolation, data export/delete (R-21, R-23)

MVP note on Trust depth

The MVP can ship with a single-item trust rating plus inferred trust from open text; the full multi-item driver battery (reliability/integrity/benevolence/security) can be a fast-follow if it risks the MVP timeline. See Q-14.

Phase 1 — Competitive insight & distribution

- Competitor configuration and benchmarking (own vs competitors) on ratings, Brand Love, Trust, sentiment, themes (R-24, R-27, O-9) — CSV / licensed-provider / API sources first.
- Live web collection (R-25, DPS-7) — enabled only after the client's legal review and sign-off.
- Per-competitor deep-dive (R-28, O-10).
- Natural-language "ask your data" query (R-17, O-6); email distribution (R-4 email); alerts (R-20); slides export.

Phase 2 — Connectors & scale

- Additional sanctioned connectors; larger-volume background processing (NFR-5); multi-language expansion (Q-4); team collaboration.

Phase 3 — Advanced insight

- Driver analysis — what most influences Brand Love and Trust; the love-type distinction (passion vs. intimacy/loyalty, per Nobre 2011) and a resilience measure (does sentiment/repurchase hold after a negative event); benchmarking vs prior periods and optional, opt-in, anonymized peer norms; a research assistant that proposes surveys from the gaps it finds.

SECTION 14

Open questions

Decisions we need from you before proceeding to CHALLENGE (Phase 2).

- **Q-1** Product name — is "AAA Insights" the intended name or a working title?
- **Q-2** Build vs. buy the survey engine — build collection natively, or wrap an existing component?
- **Q-3** AI models / hosting — any constraints on model provider or data residency?

- **Q-4** Language coverage for v1 — English-only, or which additional languages at launch?
- **Q-5** Compliance targets — which regimes must v1 satisfy (GDPR, CCPA, SOC 2 roadmap)?
- **Q-6** MVP volume target — what item volume must the MVP handle comfortably?
- **Q-7** Commercial model — per-seat, per-response, per-survey, per-competitor, or tiered?
- **Q-8** First design partner — a specific first customer/segment to anchor the MVP?
- **Q-9** Brand-love framework — resolved: grounded in Batra/Ahuvia/Bagozzi (2012). Confirm "Ambiguity" as the middle label (vs. "Ambivalence"/"Neutral").
- **Q-10** Competitors & sources — which competitors to track first, and which review sources matter most?
- **Q-11** Collection method & legal sign-off — confirm the API / licensed-provider / public-web mix, and who provides the legal review that gates live web collection (DPS-7).
- **Q-12** Rating normalization — confirm the common scale (proposed 1–5) and the Brand Love Index formula (proposed: $\%(\text{Love} + \text{Like}) - \%(\text{Dislike} + \text{Hate})$).
- **Q-13** Collected-data retention — how long should collected competitor/web data be retained?
- **Q-14** Trust depth for v1 — single-item trust rating plus inference, or the full multi-item driver battery? (Recommendation: single-item + inference in the MVP; driver battery as a fast-follow.)
- **Q-15** Trust driver taxonomy — confirm reliability / integrity / benevolence / security-privacy as the drivers, and the Trust Index formula (proposed net: $\%\text{positive-trust} - \%\text{negative-trust}$ on the normalized scale).

APPROVAL

The Phase-1 gate

Approved by _____ Date _____

Prepared by Active AI Advisors under the Grounded AI™ methodology. Every phase produces an artifact you own.

Active AI Advisors · paul@activeaiadvisors.com