



PRODUCT REQUIREMENTS DOCUMENT

AAA Insights

A feedback-to-insight service: collect surveys, reviews, ratings, and comments — then let AI turn them into decisions.

Version: v1 – Draft (awaiting client approval)
Date: 2026-07-13
Prepared by: Active AI Advisors
Prepared for: Paul Jamieson
Methodology: Grounded AI™ – Phase 1 (SPECIFY)

Applied AI for operators. · activeaiadvisors.com · paul@activeaiadvisors.com

CONTENTS

- 01 Purpose
- 02 Target users & buyer
- 03 Product principles
- 04 Inputs
- 05 Outputs
- 06 Requirements (behaviors)
- 07 Invariants
- 08 Edge cases
- 09 Exclusions
- 10 Data, privacy & security
- 11 Non-functional requirements
- 12 Scope & phased roadmap
- 13 Open questions
- Approval

SECTION 01

Purpose

What the software is for — and who it serves.

SECTION 02

Target users & buyer

Two audiences, one product: credible analysis for the business, effortless response for the customer.

User role	What they need from AAA Insights
Owner / founder / GM (buyer)	A fast, trustworthy read on what customers think and what to fix first — without hiring a research team.
Marketing / brand lead	Themes, sentiment, and verbatim quotes to guide messaging, positioning, and product decisions.
CX / customer-success manager	Early warning on dissatisfaction, drivers of churn, and a way to close the loop with unhappy customers.
Product manager	Prioritized signal on what customers want changed, backed by evidence and volume.
Respondent (the customer)	A short, respectful, mobile-friendly way to give feedback — including a conversational option that feels heard.

SECTION 03

Product principles

Five commitments that shape every design decision.

SECTION 04

Inputs

What the system receives.

- **I-1** Survey definition — built by an admin (questions, types, logic, branding) or AI-generated from a plain-language goal. Stored as structured survey config.
- **I-2** Survey response — a respondent's answers: ratings, multiple-choice, and open-text comments. May be partial if the respondent abandons midway.
- **I-3** Conversational response — free-text turns in an AI-led adaptive interview, captured as a threaded transcript of probes and replies.
- **I-4** Imported feedback — external reviews, ratings, and comments via CSV upload (MVP) and, later, connectors. Each carries source, timestamp, rating, and text.
- **I-5** Respondent metadata — optional, non-identifying attributes (segment, channel, product line, region, language) used for filtering and comparison.
- **I-6** Distribution request — an admin action to publish a survey via shareable link, email invitation list, or embeddable widget.
- **I-7** Analysis query — a natural-language question asked of the feedback data ("why are enterprise customers unhappy this quarter?").

SECTION 05

Outputs

What the system produces.

- **O-1** Collected response store — every response (survey, conversational, imported) persisted under a common schema, attributable to its source and campaign.
- **O-2** Theme analysis — AI-generated themes for a chosen set of responses, each with a label, count/percentage, representative verbatim quotes, and trend vs. prior period.
- **O-3** Sentiment analysis — positive / neutral / negative with intensity, at the individual-response level and aggregated by theme, segment, source, and time.
- **O-4** Ratings & metrics — average ratings, distribution, response and completion rates, and standard indices (NPS, CSAT) where the questions qualify.
- **O-5** Insight report — a plain-language narrative of what's happening, why, and what to do, with a short list of recommended actions ranked by evidence and impact. Exportable.
- **O-6** Answer to an analysis query — a natural-language response grounded strictly in the company's own data, citing the specific responses used.
- **O-7** Alerts — notifications when a monitored signal crosses a threshold (a negative-sentiment spike, an emerging theme, a rating floor).
- **O-8** Respondent-facing survey — the rendered survey or conversational interview the customer actually sees and completes.

SECTION 06

Requirements (behaviors)

Each is a single, testable behavior. IDs are cited by the Phase-3 test suite.

Survey creation & distribution

- **R-1** An admin can create a survey with at least these question types: single-select, multi-select, rating/scale (configurable range), and open text.
- **R-2** An admin can generate a draft survey from a plain-language objective; the system returns an editable survey to modify or discard before sending.
- **R-3** A survey supports conditional logic: a question or branch is shown or skipped based on a previous answer.
- **R-4** An admin can distribute by (a) shareable public link, (b) email invitation list, and (c) embeddable widget. (Email is a Phase-1 roadmap item; link and widget are MVP.)
- **R-5** A respondent can complete a survey on mobile and desktop without creating an account.
- **R-6** The system records partial responses and marks them incomplete rather than discarding them.

Conversational (AI-led) surveys

- **R-7** An admin can enable a "conversational" mode in which an AI interviewer asks questions in natural language and asks relevant follow-up probes based on the answers.
- **R-8** The interviewer stays on the admin-defined objective and does not ask questions outside the configured topic scope.
- **R-9** The interviewer ends after a configurable maximum number of exchanges or when the objective is met, whichever comes first.
- **R-10** Every conversational transcript is stored and is analyzable by the same theme/sentiment engine as structured responses.

Unified feedback hub

- **R-11** An admin can import external feedback via CSV, mapping columns to the unified schema (source, date, rating, text, segment).
- **R-12** All feedback — survey, conversational, imported — is queryable together, filterable by source, campaign, date range, segment, rating, and sentiment.
- **R-13** The system de-duplicates obviously identical imported responses (same source + text + timestamp) so they are not double-counted.

AI analysis & insight

- **R-14** For any filtered set of responses, the system produces theme analysis with counts and representative quotes.

- **R-15** The system assigns sentiment to each open-text response and aggregates it across the chosen dimensions.
- **R-16** For every theme and sentiment aggregate, the user can open the exact underlying responses that produced it (traceability — INV-3).
- **R-17** An admin can ask a natural-language question of their data and receive an answer grounded only in their own responses, with citations.
- **R-18** The system generates an insight report for a chosen scope: narrative, key metrics, top themes, sentiment summary, quotes, and ranked recommended actions.
- **R-19** The system computes response rate, completion rate, average rating, rating distribution, and at least NPS and CSAT when the questions qualify.
- **R-20** An admin can configure an alert on a signal (negative-sentiment share, average rating, new theme) and be notified when the threshold is crossed.

Accounts, roles & export

- **R-21** At least two roles: an Owner/Admin who manages account, users, and billing, and a Member who builds surveys and views analysis but not billing/users.
- **R-22** An admin can export any analysis view and the insight report in a portable format (CSV for raw data; PDF and/or slides for the report).
- **R-23** An admin can delete a survey, campaign, or individual response, and can export or delete all data for the account.

SECTION 07

Invariants

Conditions that must hold in every state of the system.

- **INV-1** A response is always attributable to exactly one source (link, campaign, conversational session, or named import batch). No response exists without a known origin.
- **INV-2** Response counts are never negative and never exceed responses actually collected; percentages are computed from the current filter and each response is counted once per dimension.
- **INV-3** Every AI-generated theme, aggregate, insight, recommendation, and query answer is traceable to the specific underlying responses. The UI never shows an AI claim without a path to its verbatim data.
- **INV-4** The system never displays a quote or statistic absent from the collected data. AI output summarizes only what exists; it does not invent respondents, quotes, or numbers.
- **INV-5** A respondent can complete a survey without providing personally identifying information; identity is never required to submit feedback.

- **INV-6** Each account's data is isolated: no account can see, query, or analyze another's responses. AI analysis is grounded only in that account's data.
- **INV-7** Deleting a response, campaign, or account removes it from all future analysis and aggregates; deleted data never reappears in any report.
- **INV-8** Respondent PII, when present, is stored securely and never written to logs or included in AI outputs unless the admin explicitly opts in.

SECTION 08

Edge cases

The unusual conditions the system must survive — defined now, tested in Phase 3.

- **E-1** Empty or tiny sample — the system states the sample is too small for reliable themes/sentiment rather than fabricating confident findings.
- **E-2** Abandoned survey — the partial is retained, counted as incomplete in completion math, and flagged in analysis.
- **E-3** Junk / abusive open text — captured, flagged, excluded from theme/sentiment by default, and reviewable by the admin.
- **E-4** Non-English / mixed-language responses — language is detected and analyzed accordingly (supported languages are an open question — Q-4).
- **E-5** Malformed CSV import — rejected or partially accepted with a per-row error report; no partial garbage is silently stored.
- **E-6** Duplicate / repeated submissions — de-duplication (R-13) applies, plus link-level controls (e.g., one response per link).
- **E-7** Conversational interview goes off the rails — the interviewer stays in scope, ignores injected instructions, and ends gracefully.
- **E-8** Ambiguous analysis query — the system answers only what the data supports and states what it cannot determine, rather than guessing.
- **E-9** Sudden volume spike — ingestion and analysis stay correct; real-time views may lag but never lose or double-count data.
- **E-10** Sensitive disclosure — PII or safety concerns are stored under INV-8 and surfaced to the admin without broadcasting into shared outputs.
- **E-11** Concurrent editing — two admins editing one survey do not silently overwrite each other without warning.

SECTION 09

Exclusions

Explicitly out of scope for v1 — the unwritten assumptions, made written.

- **X-1** Not a CRM or customer-record system; it stores feedback, not a system of record for customer accounts.
- **X-2** No two-way case management / ticketing; closing the loop in v1 is export/notification, not an in-app resolution workflow.
- **X-3** No scraping of review sites without authorization; external feedback enters via CSV in the MVP and sanctioned connectors later.
- **X-4** No prediction of future business outcomes (churn, revenue) in v1; it analyzes feedback that exists.
- **X-5** No public respondent panel / audience marketplace; the company brings its own audience.
- **X-6** Does not replace human research judgment; it is decision support, with every AI output presented as such alongside its evidence.

SECTION 10

Data, privacy & security

The handling bar this product is held to — enforced in IMPLEMENT and VERIFY.

- **DPS-1** The system is a data processor for the account's feedback; the company is the controller. Roles, retention, and deletion honor that relationship.
- **DPS-2** Respondent PII is minimized, encrypted at rest and in transit, access-controlled by role, and excluded from logs and AI outputs by default.
- **DPS-3** Data-subject requests are supported: export and deletion of an individual respondent's data on request.
- **DPS-4** AI processing is isolated to the account; account data is not used to train shared/base models without explicit opt-in.
- **DPS-5** A configurable retention period is supported; data past retention is deleted and drops out of analysis.
- **DPS-6** The full security checklist (input validation, authz/authn, secret handling) is applied during build and verification.

SECTION 11

Non-functional requirements

Speed, reliability, access, and honesty — stated as testable budgets.

- **NFR-1** Performance — a survey page loads in under 2 seconds on a typical mobile connection; a submission acknowledges in under 1 second.
- **NFR-2** Analysis latency — theme/sentiment analysis up to the MVP size (target ~5,000 responses, Q-6) returns within a stated budget; larger sets run as background jobs with progress shown.
- **NFR-3** Availability — the collection endpoint targets a stated uptime so responses are never lost during a live window; ingestion is durable even if analysis is degraded.
- **NFR-4** Accessibility — respondent-facing surveys meet WCAG 2.1 AA.
- **NFR-5** Scalability path — the architecture grows from SMB volumes toward larger sets without a rewrite; the MVP need not deliver enterprise scale but must not preclude it.
- **NFR-6** Explainability — every AI output carries an indication of the data behind it and, where meaningful, a confidence signal.

SECTION 12

Scope & phased roadmap

Only Phase 0 is committed by this document. Each later phase re-enters the pipeline before it is built.

Phase 0 — MVP (the focused, buildable core)

The smallest release that delivers the differentiated promise end-to-end for one SMB account.

- Survey builder: four core question types + conditional logic (R-1, R-3) • AI-drafted survey from a goal (R-2)
- Distribution by link + embeddable widget (R-4) • Mobile, account-free capture incl. partials (R-5, R-6)
- Conversational AI survey with scope control + turn limits (R-7–R-10) • CSV import into the unified hub (R-11, R-13)
- Unified analysis with filtering (R-12) • AI theme + sentiment with traceability (R-14–R-16)
- Core metrics incl. NPS/CSAT (R-19) • Insight report with ranked actions, exportable (R-18, R-22)
- Two roles, account isolation, data export/delete (R-21, R-23, INV-6, INV-7)

Phase 1 — Depth & distribution

- Natural-language "ask your data" query with citations (R-17, O-6).
- Email invitation distribution and reminders (R-4 email).
- Configurable alerts on sentiment/rating/theme signals (R-20, O-7).
- Segment comparison and trend-over-time views; slides export.

Phase 2 — Connectors & scale

- Sanctioned connectors for external review/rating sources and helpdesk exports (replacing CSV-only).
- Larger-volume analysis with background processing at scale (NFR-5).
- Multi-language analysis expansion (per Q-4).
- Team collaboration: shared workspaces, comments on insights, saved views.

Phase 3 — Advanced insight

- Driver analysis — what most influences satisfaction or rating.
- Benchmarking against prior periods and optional, opt-in, anonymized peer norms.
- A research assistant that proposes surveys and follow-up studies from the gaps it finds.

SECTION 13

Open questions

Decisions we need from you before proceeding to CHALLENGE (Phase 2).

- **Q-1** Product name — is "AAA Insights" the intended name or a working title?
- **Q-2** Build vs. buy the survey engine — build collection natively, or wrap an existing component and focus the build on the AI analysis layer? Materially affects timeline and cost.
- **Q-3** AI models / hosting — any constraints on model provider or data residency to design to from day one?
- **Q-4** Language coverage for v1 — English-only, or which additional languages at launch?
- **Q-5** Compliance targets — which regimes must v1 satisfy (GDPR, CCPA, SOC 2 roadmap)?
- **Q-6** MVP volume target — what response volume must the MVP handle comfortably? Sets NFR-2 budgets and the storage/analysis approach.
- **Q-7** Commercial model — per-seat, per-response, per-survey, or tiered? Shapes account, quota, and billing needs.
- **Q-8** First design partner — is there a specific first customer/segment to anchor the MVP to?

- **Q-9** Brand-love research inputs — should the insight engine build in a specific research framework (e.g., brand-love / emotional-attachment constructs) as a lens, or stay framework-neutral in v1?

APPROVAL

The Phase-1 gate

Approved by _____ Date _____

Prepared by Active AI Advisors under the Grounded AI™ methodology. Every phase produces an artifact you own.

Active AI Advisors · paul@activeaiadvisors.com